Introduction to machine learning

RNN

Introducción al aprendizaje de la máquina

DeepMind

Self-driving cars

Cloud optimization

Healthcare
&
Medical images analysis

GAN

This is not a real cup of coffee

~~12/03/2020~~

~~30/04/2020~~

~~14/05/2020~~

May 20th 2020

Luc Lesoil

Presentation

I. Supervised learning

Break – A video to introduce DiverSE

II. Unsupervised learning

III. Reinforcement learning

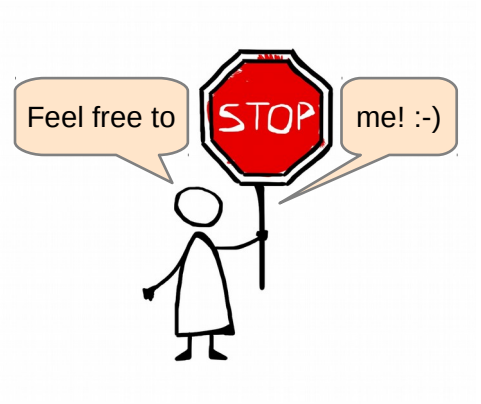

Introduction

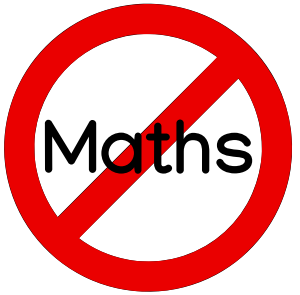

- High-level presentation
- Non-exhaustive list
- Concrete cases

https://github.com/llesoil/ML_example

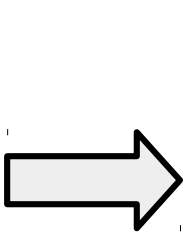

I- Supervised learning

X : images of numbers

y : numbers

3 5 7
4 0
1

X : Explaining variables

y : Variable to predict, labels known

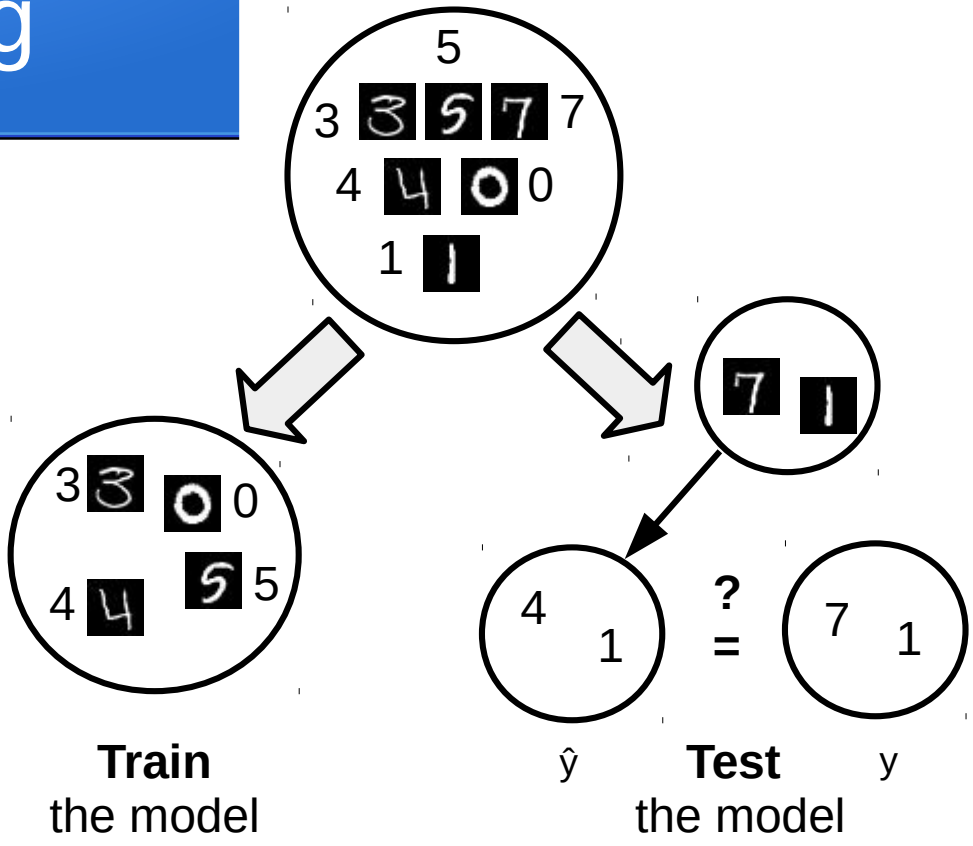

Supervised learning = Use X to predict y

Loss function

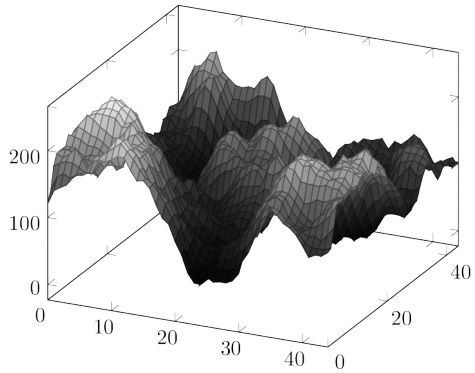

5

Today's Tux happiness : 3.2

2.5

Prediction 1 : 3

0

Prediction 2 : 0.5

Compare quality of predictions?

Loss function

Minimize the loss function

=

better predictions

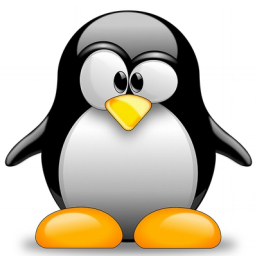

Tux

https://towardsdatascience.com/common-loss-functions-in-machine-learning-46af0ffc4d23

https://algorithmia.com/blog/introduction-to-loss-functions

Examples	
MAE	MAPE
Minkowski	Hinge
MSE	Cross-entropy

Type of

Supervised learning

Classification

→ Group or category

Regression

→ Value

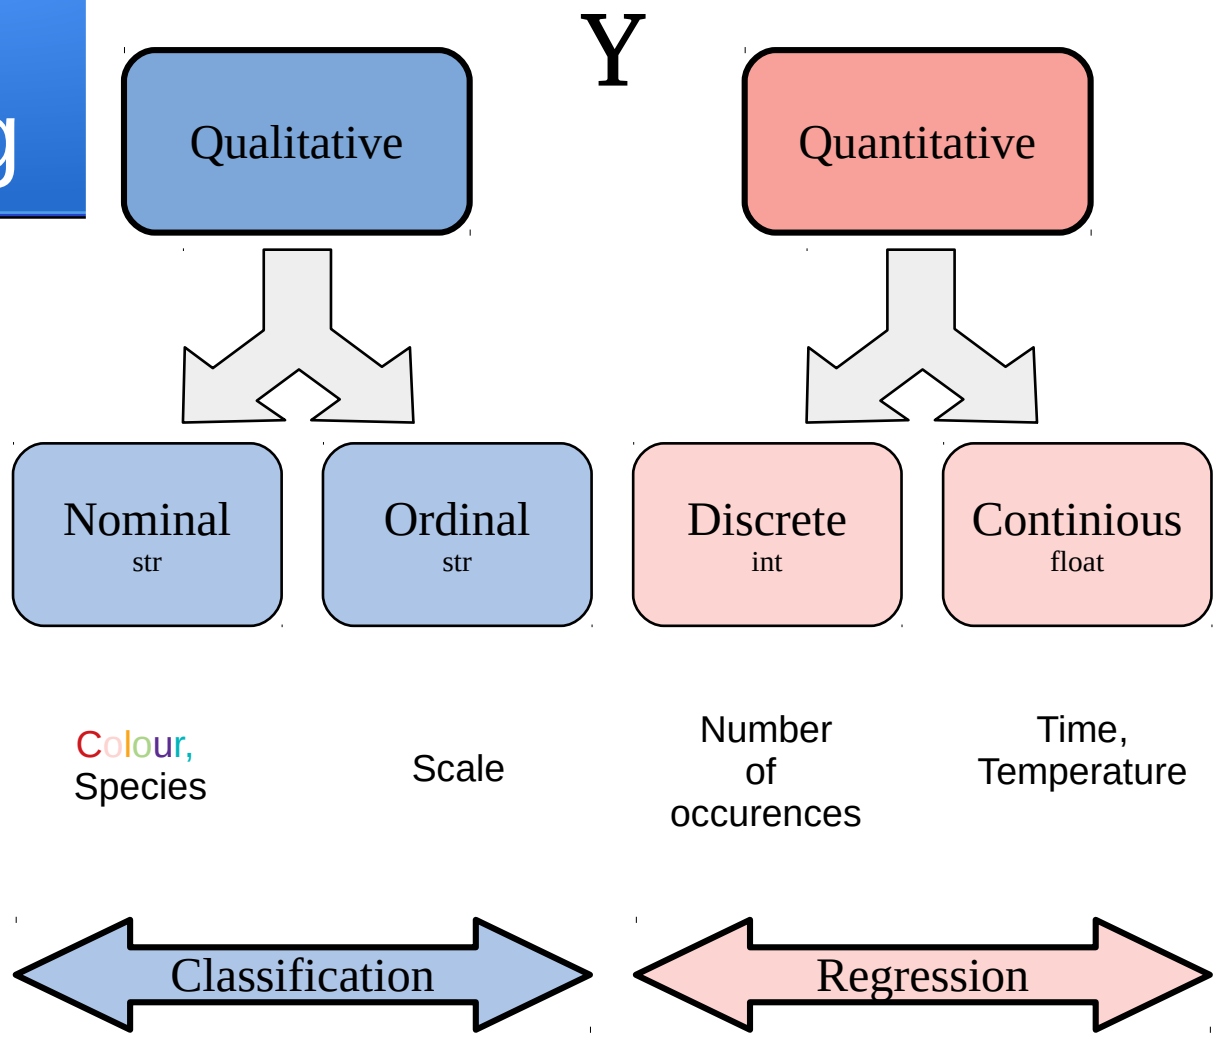

Linear regression

- Simple
- ~~Complex dataset~~
- Linear relationship

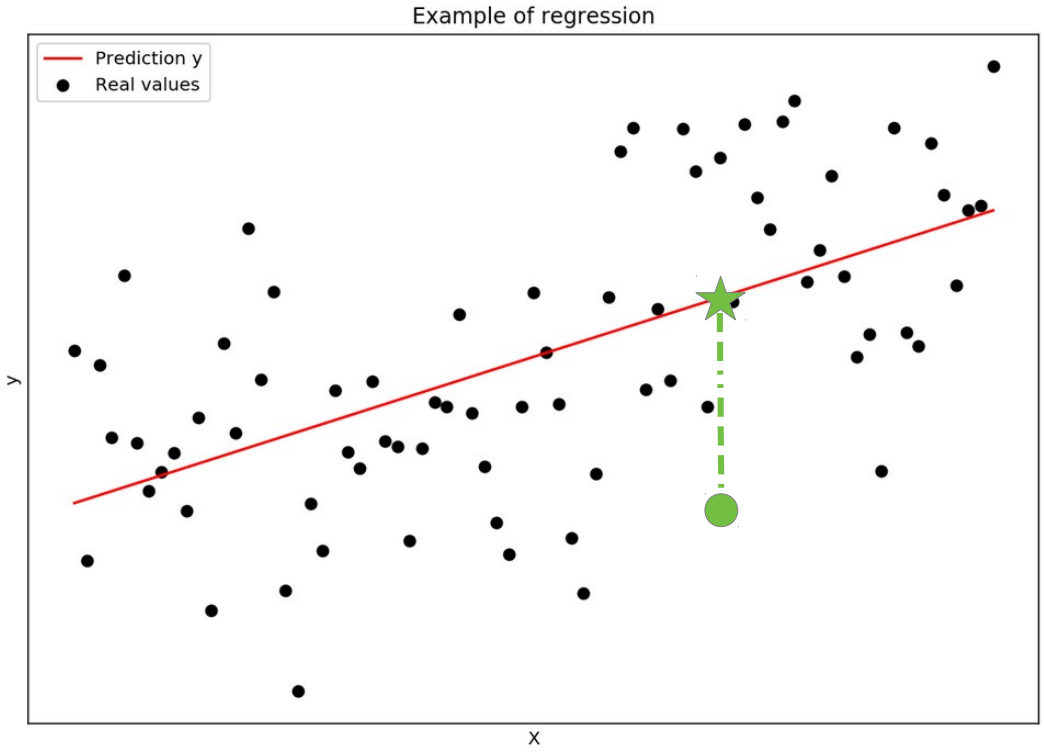

Fit the scatterplot with the red line

★ is the prediction of x = ●

Decision Tree (CART)

- Extract rules
- Simple to parameter
- Learning unit for many algorithms

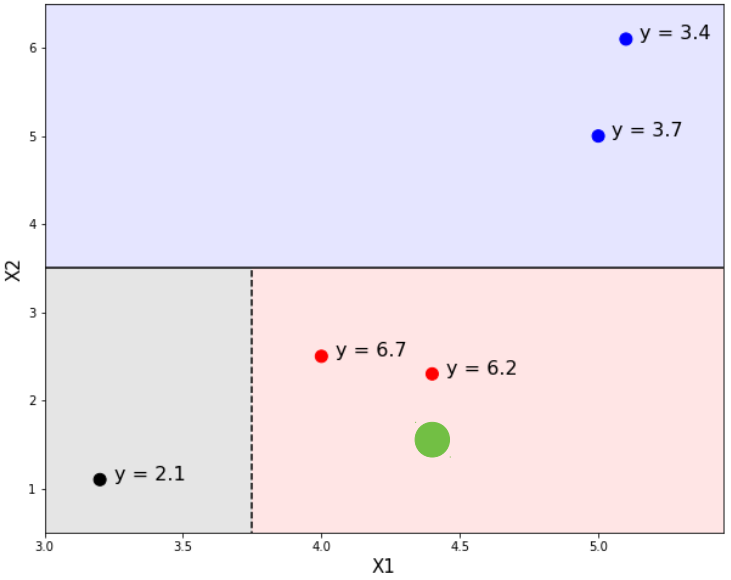

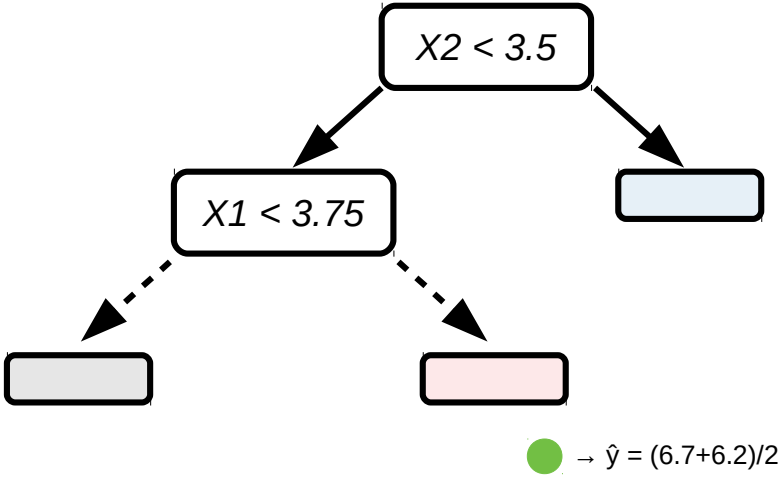

Classification Regression

Random Forest

- Bagging → robustness
- Metrics
- Good compromise

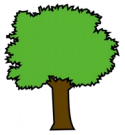

= 1 Decision Tree

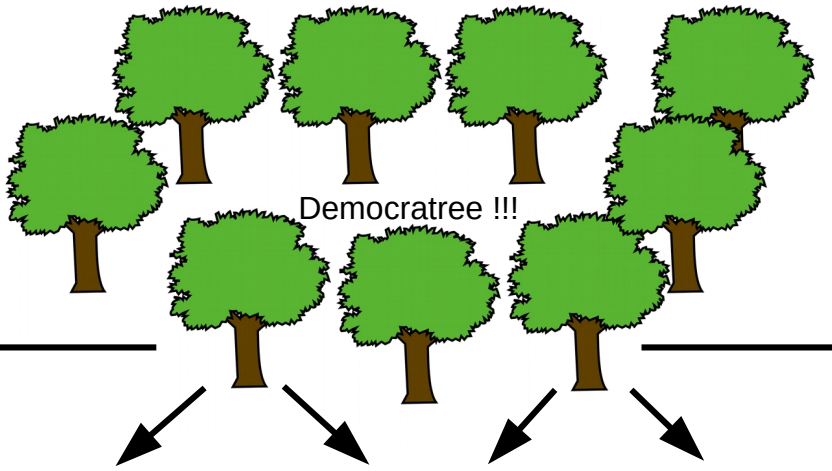

X_1	X_2	X_3
1	7	6
4	3	0
8	1	2
2	4	3

Classification Regression

Boosting tree

- Complex dataset
- Many hyperparameters
- XGBoost: the algorithm that wins every competition

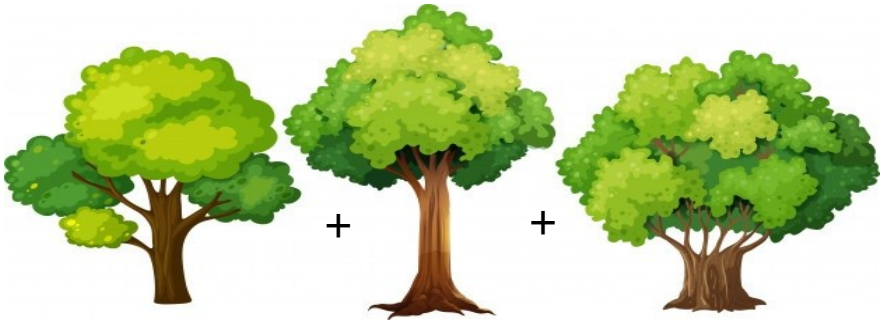

Update the trees based on previous results

AdaBoost

XGBoost

Classification Regression

Neural networks

- ~~Simple dataset~~
- Many hyperparameters
- Black box

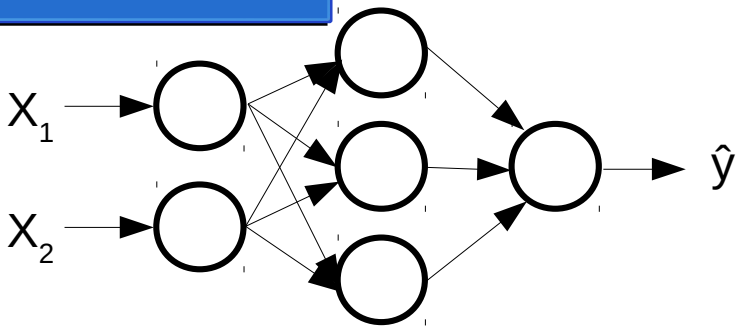

Feedforward neural network

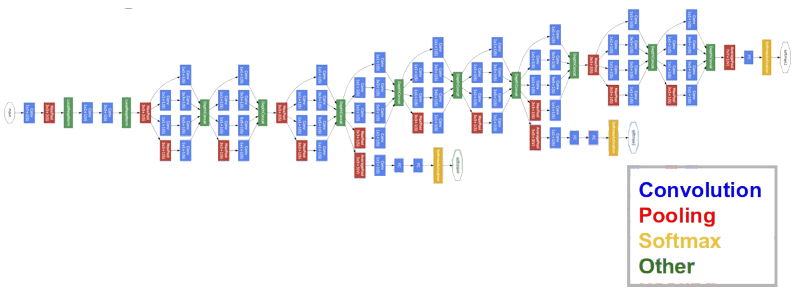

Classification Regression

Others

- Quantile/Polynomial/Piecewise regression, Ridge, ElasticNet, LASSO to select explaining variables

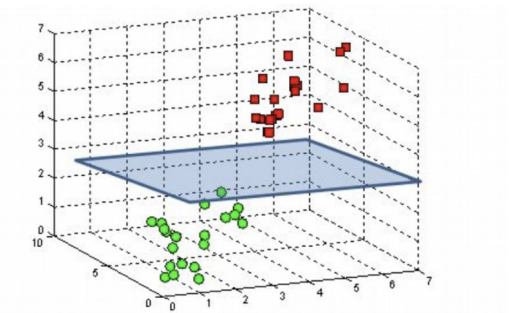

- Support Vector Machine : SVC or SVR

- Time series predictions : (S)AR(I)MA, RNN

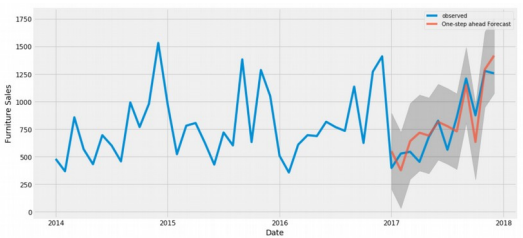

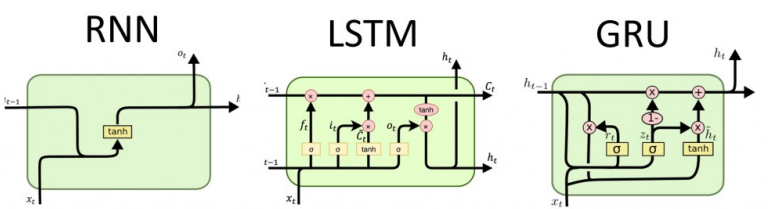

Break – Video

- Introduce the team
- YouTube channel
- 30 seconds

Can we use your photo?

You want to add a software in the list?

You have a good microphone and-or a beautiful voice?

You want to share a demo?

You want to write the description of the team?

II- Unsupervised learning

- Clustering
- Association
- Anomaly detection

Supervised Learning

- Step: 1 X → y

  Teacher: Does picture 1 show a car or a flower?

  Learner: A flower.

  Teacher: No, it's a car.

Step: 2

  Teacher: Does picture 2 show a car or a flower?

  Learner: A car.

  Teacher: Yes, it's a car.

Step: 3 ....

	X	y
Supervised	✓	✓
Unsupervised	✓	✗

Kmeans

- Simple clustering
- Fast
- Few parameters

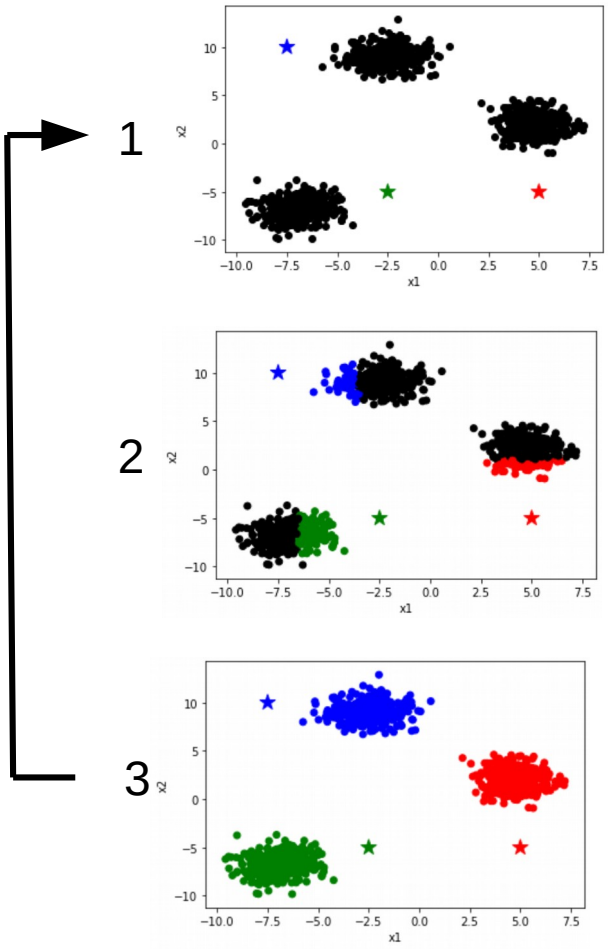

Gaussian Mixture Model

- Gaussian distribution
- Estimation of K
- Scale well - fast

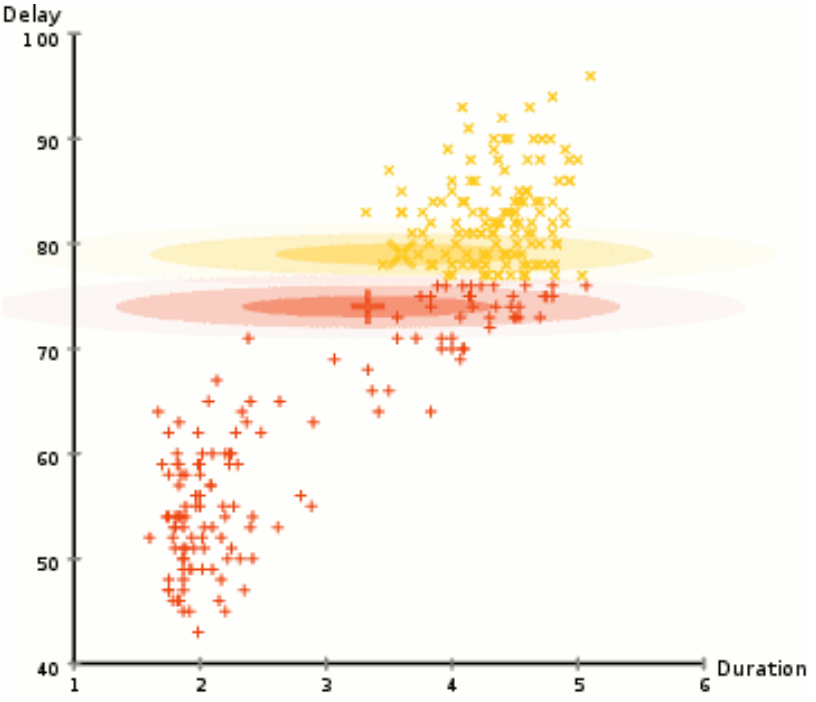

K-Nearest neighbors

- Used in recommendation systems
- Supervised
- “You are the average of the five people you spend the most time with”

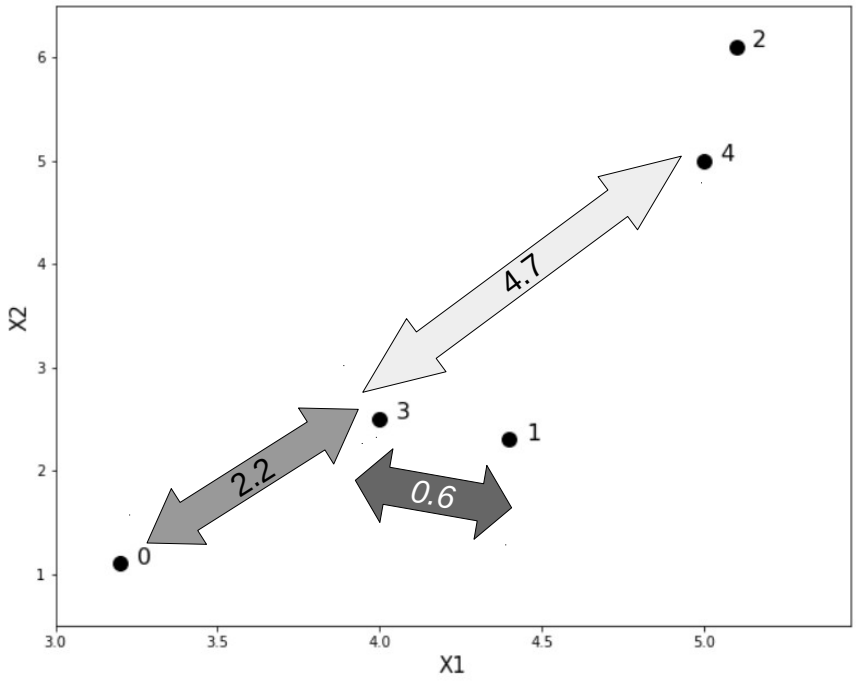

1 is the nearest neighbor of 3
0 is the second nearest

Hierarchical clustering

- Quadratic O(n²)
- Not designed for big dataset
- Full description of relationships

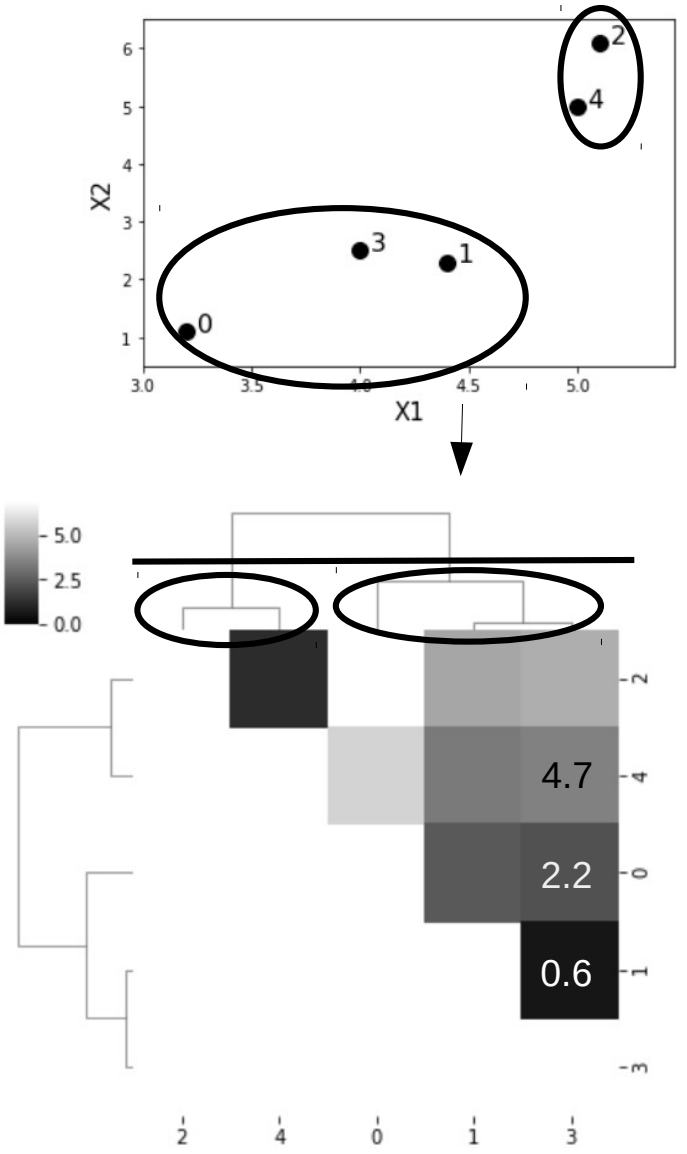

Hierarchical clustering (2)

- Quadratic O(n²)
- Not designed for big dataset
- Full description of relationships

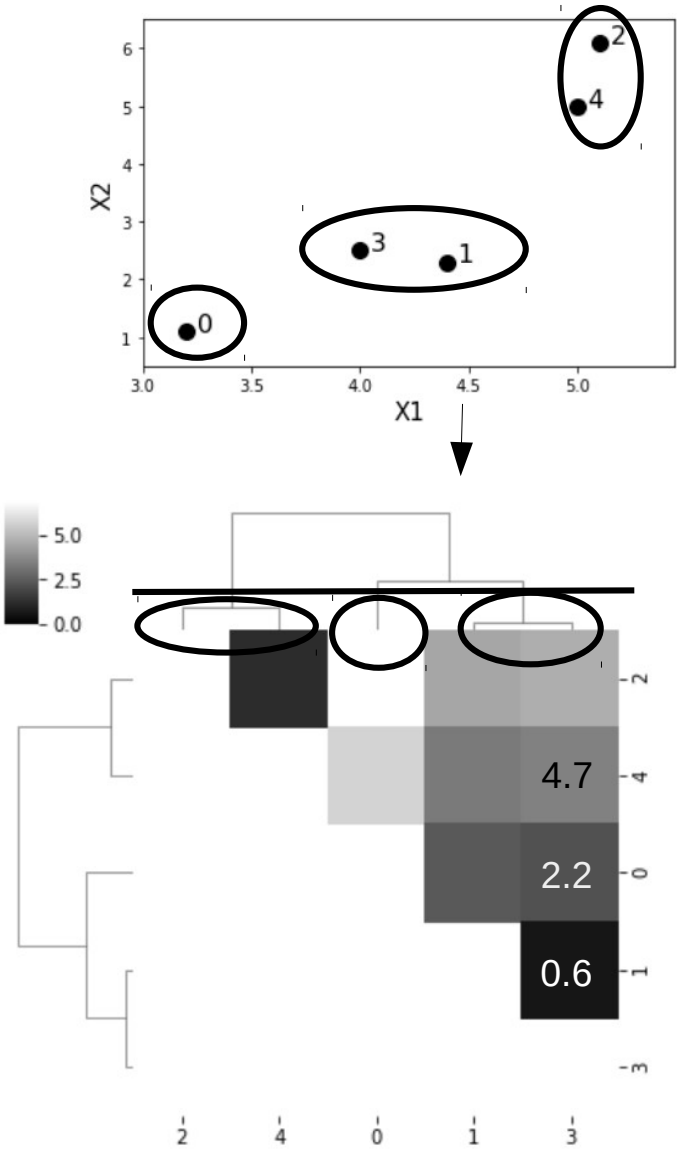

III- Reinforcement Learning

Based on behavioral psychology

Realistic learning

Google - energy consumption -15%

Traffic control

Notions

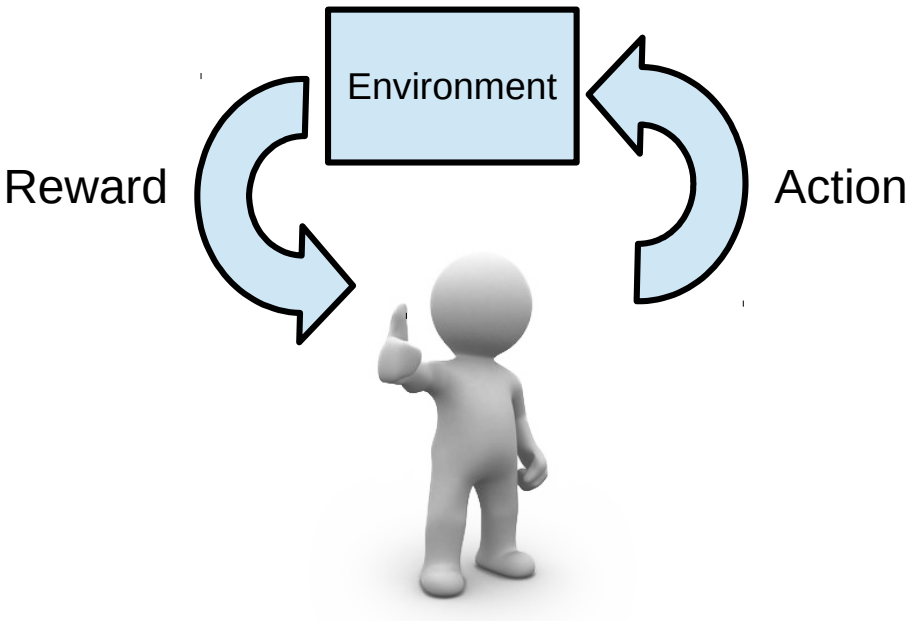

Agent, State

Policy

- Goal
- (State, Action) → Reward
- Q-Table

Reinforcement Learning

Step: 1

World: You are in state 9. Choose action A or C.

Learner: Action A.

World: Your reward is 100.

Step: 2

World: You are in state 32. Choose action B or E.

Learner: Action B.

World: Your reward is 50.

Step: 3 ....

Example

References

- Mnist dataset, scikit-learn documentation
- https://internetofbusiness.com/google-using-deepmind-ai-to-reduce-energy-consumption-by-30/
- https://www.slideshare.net/cprakash2011/reinforcement-learning-40052403/5
- https://brilliant.org/wiki/gaussian-mixture-model/
- https://perfectial.com/blog/q-learning-applications/
- MARIQ : https://www.youtube.com/watch?v=CacRZmjDIr4
- The egg, Andy Weir